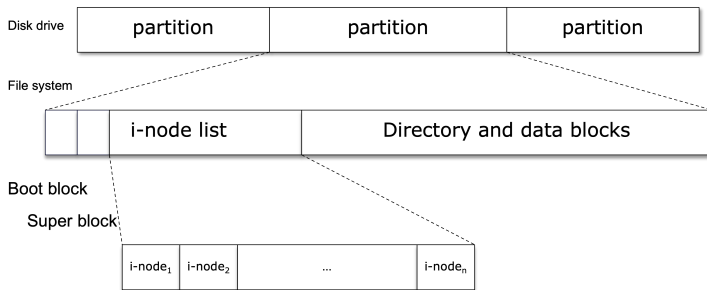


# Linux File Stats

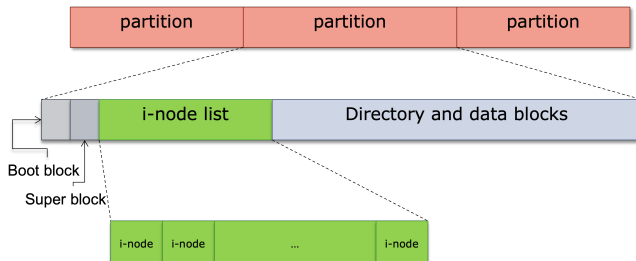
# File Partition

- Various implementations of the UNIX/Linux file system are in use today.
- We can think of a disk drive being divided into one or more partitions. Each partition can contain a file system.



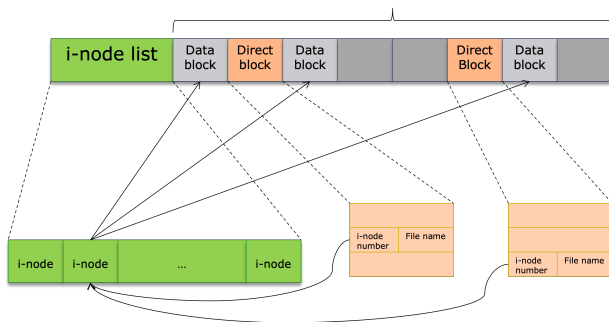
# I Node List

- In Linux, almost everything is represented as a file.
- Each disk partition can be divided into disk blocks of smaller size. A block size varies in different systems.
- The Linux file system uses Index-Node (aka. I-Node) implementation method to keep track of how many blocks, and which blocks are used for a file.
- The superblock essentially records a file system's characteristics such as block size, other block properties, sizes of block groups and location of inode tables.



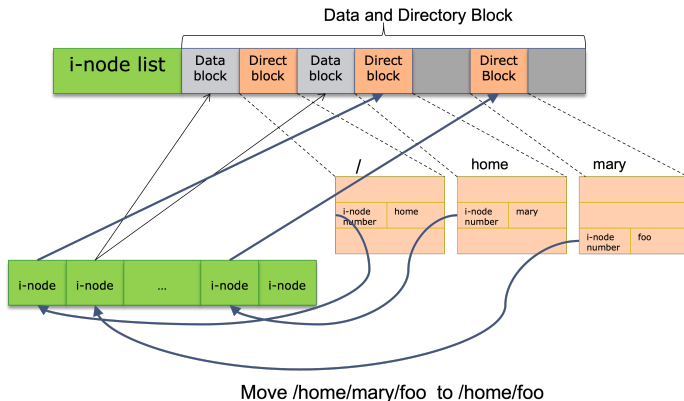
# Directories and Data Blocks

- A Linux directory is a special file that acts as a container for other files and subdirectories.
- A directory entry stores information needed to locate the file disk blocks or subdirectory. Specifically, it stores i-node number, and filename/subdirectory name.
- The i-node stores file related information: file type, access permission, etc...
- Every i-node also stores the number of links from directory entries pointing to the i-node.
- A file is deleted when the number of links is 0.



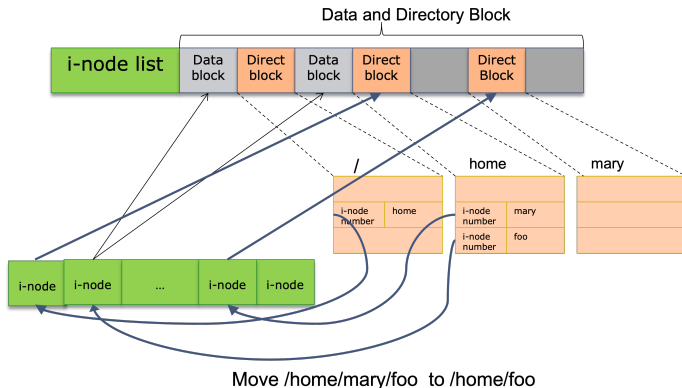
# Example of File Move

- When move a file from one directory to another directory, it only needs to change directory entry point.



# Example of File Move (Cont')

- When move a file from one directory to another directory, it only needs to change directory entry point.



# stat(), fstat(), lstat() System Calls for File Attributes

- Information/attributes related to a file can be obtained through `stat()`, `fstat()`, or `lstat()` system call.
- The `stat()` returns the structural information on a given file.
- The `fstat()` obtains attributes on a file with an open descriptor.
- The `lstat()` function is similar to `stat`, but when the named file is a symbolic link, `lstat()` returns information about the symbolic link.

Prototype of the 3 system calls:

```
#include <sys/types.h>
#include<sys/stat.h>

// returns 0 if success, -1 on error.
int stat(const char *fname, struct stat *buf);
int fstat(int fd, struct stat *buf);
int lstat(const char *fname , struct stat *buf);
```

# Definition of stat() System Call

```
struct stat {
    mode_t st_mode;    /*file type & mode (permissions) */
    ino_t st_ino;      /* i-node number */
    dev_t st_dev;      /* device number (file system) */
    dev_t st_rdev;      /* device number for special files */
    nlink_t st_nlink;  /* number of links */
    uid_t st_uid;      /* user ID of owner */
    gid_t st_gid;      /* group ID of owner */
    off_t st_size;     /* size in bytes, for regular files */
    time_t st_atime;    /* time of last access */
    time_t st_mtime;    /* time of last modification */
    time_t st_ctime;    /* time of last file status change */
    blksize_t st_blksize; /* best I/O block size */
    blkcnt_t st_blocks; /*number of 512 byte blocks allocated */
    mode_t st_attr;     /* The DOS-style attributes for this file */
};
```

The above special data types are all aliases for integer types (signed or unsigned):

[https://www.gnu.org/software/libc/manual/html\\_node/Attribute-Meanings.html](https://www.gnu.org/software/libc/manual/html_node/Attribute-Meanings.html)

or type `man 2 stat`



# File Types

**Regular file** contains data of some form (text or binary)

**Directory file** contains the name of other files and pointers to the information on these file.  
Only kernel can write to the directory file

**Character special file** is a type of file providing unbuffered I/O access to variable-sized units to devices.

**Block special file** is a type of file used for disk devices.

**FIFO** is a type of file used for inter-process communication.

**Socket** is a type of file used for used for network communication.

**Symbolic link** is a type of file that points to another file.

# How to Determine File Types

- The type of a file is encoded in the `st_mode` member of the `stat` structure.
- We can determine the file type by using macros in `<sys/stat.h>`.

```
S_ISREG(st_mode): Regular file
S_ISDIR(st_mode): Directory
S_ISCHR(st_mode): Character special
S_ISBLK(st_mode): Block special
S_ISFIFO(st_mode): FIFO
S_ISLNK(st_mode): Symbolic link
S_ISSOCK(st_mode): Socket
```

# Code Example to Display File Type

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <sys/stat.h>

void err_ret(char *str){
    printf ("%s",str);
    exit (1);
}

int main (int argc, char *argv[]){
    int i;
    struct stat sb;
    char *ptr;

    if (argc != 2)
        err_ret("Argument number error.\n");
    /* get file information */
    if (stat(argv[1], &sb) < 0)
        err_ret("lstat Error");
    if (S_ISREG(sb.st_mode)) ptr ="regular";
    else if (S_ISDIR(sb.st_mode)) ptr ="directory";
    else if (S_ISCHR(sb.st_mode)) ptr ="character special";
    else if (S_ISBLK(sb.st_mode)) ptr = "block special";
    else if (S_ISFIFO(sb.st_mode)) ptr="fifo";
    else if (S_ISLNK(sb.st_mode)) ptr = "symbolic link";
    else if (S_ISSOCK(sb.st_mode)) ptr ="socket";
    else ptr ="unknown mode ";
    printf ("%s \n", ptr);
    exit (0);
}
```

# In-depth Code Example to Display File Attributes

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <time.h> // for ctime
//Check a file type and other attributes
int main(int argc, char *argv[]){
    struct stat sb;
    if (argc != 2) {
        fprintf(stderr, "Usage: %s <pathname>\n", argv[0]);
        exit(1);
    }
    // get a file information and save in stat structure type
    if (stat(argv[1], &sb) == -1) {
        perror("stat");
        exit(1);
    }
    // check and print file type here ...
    // print other file attributes information
    printf("I-node number: %ld\n", (long) sb.st_ino);
    printf("Mode: %lo (octal)\n", (unsigned long) sb.st_mode);
    printf("Link count: %ld\n", (long) sb.st_nlink);
    printf("Ownership: UID=%ld GID=%ld\n",
        (long) sb.st_uid, (long) sb.st_gid);
    printf("Preferred I/O block size: %ld bytes\n", (long) sb.st_blksize);
    printf("File size: %ld bytes\n", (long) sb.st_size);
    printf("Blocks allocated: %ld\n", (long) sb.st_blocks);
    printf("Last status change: %s", ctime(&sb.st_ctime));
    printf("Last file access: %s", ctime(&sb.st_atime));
    printf("Last file modification: %s", ctime(&sb.st_mtime));
    exit(0);
}
```

## Question

What's the i-node number for root directory?

# link() System Call

- Any file can have multiple directory entries pointing to its i-node.
- The way we can create a link to an existing file is with the link system call.
- The link() system call create a new directory entry `newpath` that references the existing file `existingpath`.
- The link creation and link count increase are done automatically by kernel.
- Only a superuser process can create a new link that points to a directory. (Because link system calls can cause loops in the filesystem.)

```
#include <unistd.h>
```

```
//Return 0 on success, -1 on error.
```

```
int link (const char *existingpath, const char *newpath)
```

# link() System Call Example

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/stat.h>

void err_sys(char *str){
    printf ("%s\n",str);
    exit (1);
}

int main (int argc, char *argv[]){
    struct stat buff;
    if (argc < 2)
        err_sys ("Less than two argument Error");
    /* increase the number of link by one */

    if (stat (argv[1], &buff) < 0)
        err_sys("stat error for foo");
    printf ("link count for a file %s was %d \n",argv[1], buff.st_nlink);

    if (link (argv[1], argv[2]) <0)
        err_sys("Link Error");

    if (stat (argv[1], &buff) < 0)
        err_sys("stat error for foo");
    printf ("link count for a file %s is now %d \n",argv[1], buff.st_nlink);
    return 0;
}
```

# unlink() System Call

- unlink() system call removes an existing directory entry, and decrement the link count of the file referenced by pathname.
- To unlink a file, we must have write permission and execute permission in the directory containing the directory entry.

```
#include <unistd.h>
```

```
//Return 0 on success, -1 on error.
```

```
int unlink (const char *pathname)
```



# link() and unlink() System Call Example

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/stat.h>
void err_sys(char *str){
    printf ("%s\n",str);
    exit (1);
}

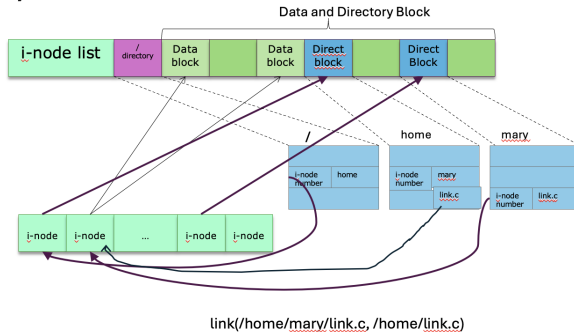
int main (int argc, char *argv[]){
    struct stat buff;
    if (argc < 2)
        err_sys ("less than two argument Error");
    if (stat (argv[1], &buff) < 0)
        err_sys("stat error for foo");
    /* increase the number of link by one */
    printf ("link count for a file %s was %d \n",argv[1], buff.st_nlink);
    if (link (argv[1], argv[2]) < 0)
        err_sys("Link Error");
    if (stat (argv[1], &buff) < 0)
        err_sys("stat error");
    printf ("link count for a file %s is now %d \n",argv[1], buff.st_nlink);
    sleep (10);
    /* unlink existing file */
    if (unlink(argv[1])<0)
        err_sys("unlink Error");

    return 0;
}
```

## Limitation of (Hard) Link

- It requires the file reside in the same file system.
- Only super user can create (hard) link to a directory.

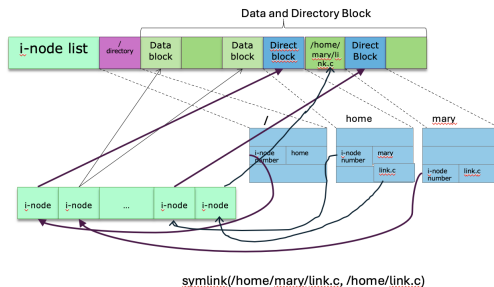
## (hard)link



# Symbolic Links

- A symbolic link is an indirect pointer to a file or directory. It stores the path of the existing file or directory.
- Unlike (hard) link, symbolic link can be used to move a file or an entire directory hierarchy to some other location on a system.
- It can also be used to execute a file.

## Symbolic link



# symlink() System Call

```
#include <unistd.h>
// return 0 if success; -1 on error.
int symlink (const char *existingpath, const char *sympath)
```

Example code to create a symbolic link to an executable file

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

void err_sys(char *str){
    printf ("%s\n",str);
    exit (1);
}

int main(int argc, char *argv[]){
    /* create a symbolic link from argv[2] to argv[1] */
    if ( symlink(argv[1], argv[2]) <0)
        err_sys("Symbolic Link Creation Error \n");

    printf("A symbolic link is created \n");

    return 0;
}
```

# Time Attributes in File Stats

Three time fields are maintained for each file.

`st_atime`

`st_mtime`

`st_ctime` – last change time of i-node status (permission, user id, number of links, etc.)

The access time and the modification time of a file can be changed with the `utime` system call.

```
//Return 0 on success, -1 on error.
```

```
#include <sys/types.h>
```

```
#include <utime.h>
```

```
int utime(const char *pathname, const struct utimbuf *times);
```

```
struct utimbuf{  
    time_t actime;  
    time_t modtime;  
}
```

# utime() System Call Example Code

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/stat.h>
#include <utime.h>
void err_sys(char *str) {
    printf ("%s\n",str);
    exit (1);
}
int main (int argc, const char *argv[]){
    struct stat statbuf0, statbuf1;
    struct utimbuf timebuf0, timebuf1;
    if (argc < 3)
        err_sys ("Argument Error");
    /* get first file information */
    if (stat(argv[1], &statbuf0) <0)
        err_sys("Stat Error");
    timebuf0.actime = statbuf0.st_atime;
    timebuf0.modtime = statbuf0.st_mtime;
    /* get second file information */
    if (stat (argv[2], &statbuf1) <0)
        err_sys("Stat Error");
    timebuf1.actime = statbuf1.st_atime;
    timebuf1.modtime = statbuf1.st_mtime;
    /* change time information between two file */
    if (utime (argv[1], &timebuf1) <0)
        err_sys("Time change Error1");
    if (utime (argv[2],&timebuf0) <0)
        err_sys("Time Change Error");
    exit (0);
}
```

# Time and Date Service in Linux

- Linux time counts the # of seconds that has passed since the Epoch:00:00:00 Jan. 1 1970.
- The time function returns the current time and date.

```
#include <time.h>
time_t time (time_t *calptr);
```

- The two functions localtime and gmtime convert a calendar time into tm structure.

```
#include <time.h>
struct tm *gmtime(const time_t *timer); /*convert to UTC */
struct tm *localtime(const time_t *timer); /*convert to local time*/

struct tm {
    int tm_sec; /* seconds [0,61] */
    int tm_min; /* minutes [0,59] */
    int tm_hour; /* hour [0,23] */
    int tm_mday; /* day of month [1,31] */
    int tm_mon; /* month of year [0,11] */
    int tm_year; /* years since 1900 */
    int tm_wday; /* day of week [0,6] (Sunday = 0) */
    int tm_yday; /* day of year [0,365] */
    int tm_isdst; /* daylight savings flag */
}
```

# Time and Date Service in Linux

- The `asctime` and `ctime` function produce the 26 byte formatted string such as  
Tue Sep 23 07:07:21 2020

```
#include <time.h>
char *asctime(const struct tm *timeptr);
char *ctime(const time_t *clock);
```

- The two functions are essentially the same, but `ctime` tend to adjusts to the time zone and daylight saving time.

---

```
/* example.c displays current date and time */
#include <stdio.h>
#include <time.h>
int main() {
    time_t t1 = time(NULL); //current calendar time in sec. since 1/1/1970
    printf("current time is: %s\n",asctime(localtime(&t1)));
    return 0;
}
```